

Demo 1

Estimating Chicago Taxi Trip Duration with the Functional API in Keras

2019

Agenda

Data Exploration

Review of dependent and independent variables

Preprocessing

Steps necessary for preparing the data for modeling

Modeling

Overview of architecture and tuning

Evaluation

Discussion of evaluation metrics and results

Improving Chicago Taxi Service by Predicting Trip Duration

Machine Learning can be used to improve the customer experience by predicting the duration of their trip on the onset of the ride

- Millions of past rides used for training and evaluating the model
- Only attributes available at the onset of the ride are used for model building
- Trip durations are predicted within 41% of the actual trip duration on 1000 unseen test rides

Data Exploration

Raw table:

`bigquery-public-data:chicago_taxi_trips.taxi_trips`

- 187,002,005 total rows
- 67.71 GB in size

Exploring Features (Independent Variables)

Potential features for modeling fall into three different categories

Predictive and Available

These features are potentially predictive of the target attribute and are available at the start of the ride

Feature	Type
trip_start_timestamp	Timestamp
pickup_census_tract	Categorical
dropoff_census_tract	Categorical
pickup_community_area	Categorical
dropoff_community_area	Categorical
pickup_latitude	Float
pickup_longitude	Float
dropoff_latitude	Float
dropoff_longitude	Float

Not Predictive or Not Available

These features are in the public Chicago Taxi dataset but are either not predictive of trip duration or not available at the start of the ride.

Feature	Type
unique_key	String
taxi_id	String
trip_end_timestamp	Timestamp
trip_miles	Float
fare	Float
tips	Float
toles	Float
extras	Float
trip_total	Float
payment_type	Categorical
company	String

Engineered

These additional fields are added to the feature set during the preprocessing phase using the available fields.

Feature	Description
distance	The distance between the pickup and dropoff lat / long values using the ST_DISTANCE() function in BigQuery
hour_start	Hour of the trip_start_timestamp field
month_start	Month of the trip_start_timestamp field
weekday	Day of week of the trip_start_timestamp field

Feature Selection

Not all available features are used for modeling

Predictive and Available

All available and predictive features are included in the model with the exception of trip_start_timestamp, as this is represented by the engineered features

Feature	Type
trip_start_timestamp	Timestamp
pickup_census_tract	Categorical
dropoff_census_tract	Categorical
pickup_community_area	Categorical
dropoff_community_area	Categorical
pickup_latitude	Float
pickup_longitude	Float
dropoff_latitude	Float
dropoff_longitude	Float

Not Predictive or Not Available

These features are not included in the model as they are either not predictive of trip duration or not available at the start of the ride.

Feature	Type
unique_key	String
taxi_id	String
trip_end_timestamp	Timestamp
trip_miles	Float
fare	Float
tips	Float
toles	Float
extras	Float
trip_total	Float
payment_type	Categorical
company	String


Engineered

All engineered features are included in the model.

Feature	Description
distance	The distance between the pickup and dropoff lat / long values using the ST_DISTANCE() function in BigQuery
hour_start	Hour of the trip_start_timestamp field
month_start	Month of the trip_start_timestamp field
weekday	Day of week of the trip_start_timestamp field

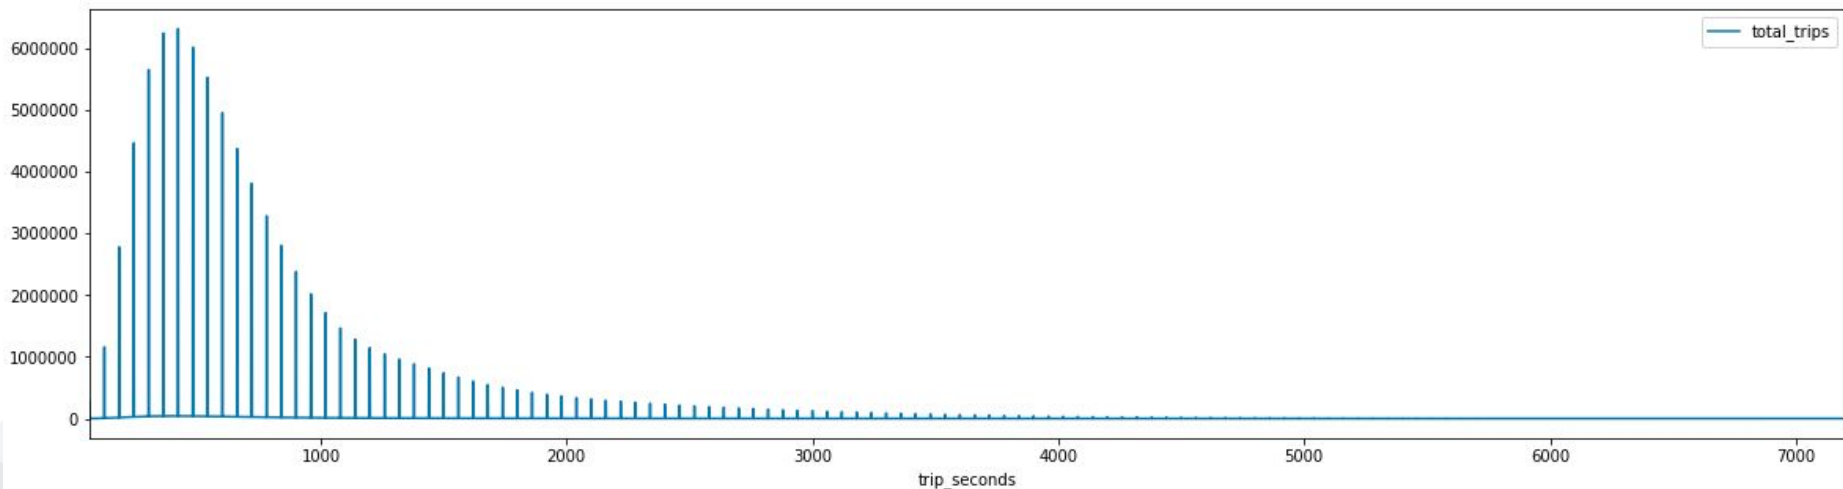
Exploring the Target Attribute (Dependent Variable): Trip Duration in Seconds

- Overall the distribution of number of trips by duration has a very long tail with some trips in the dataset lasting up to two days
- To remove outliers and improve model performance, only trips with a duration up to two hours were included for training




Exploring the Target Attribute: Log-normal Distribution

- Even after controlling for outliers, our target attribute still has a long tail, or a “log-normal” distribution
- Training a regression model where the target attribute has this type of distribution can lead to large coefficients and slow or no learning



Exploring the Target Attribute: Normalization

- To control for this, we take the natural logarithm of the trip duration as a preprocessing step prior to training. The result is that our target attribution has a much more normal distribution.
- Note this means that our predicted values will also be the natural logarithm of seconds. To convert this back to seconds, we implement a custom prediction method to return predicted trip duration in seconds by passing the initial prediction through the exponential function.



Feature Selection Summary

Continuous Features

distance

hour_start

month_start

pickup_latitude

pickup_longitude

dropoff_latitude

dropoff_longitude

Categorical Features

weekday

pickup_census_tract

dropoff_census_tract

pickup_community_area

dropoff_community_area

Target Attribute

log_trip_seconds

Preprocessing

Preprocessing occurs in two stages:

- Generating artifacts necessary for training and inference
- Preprocessing rows prior to training and / or passing for inference

Preprocessing Artifact: Training, Test, and Validation Sets

Preprocess and split raw data using BigQuery

- Transformations such as calculating logarithms and distances are implemented directly in BigQuery syntax
- Raw data is filtered based on requirements for target attribute
- Missing values are addressed
- Extracted features are split into training, test, and validation tables in BigQuery and files in Cloud Storage



train_results

113,148,861 rows (approx. 80%)

test_results

14,139,546 rows (approx. 10%)

val_results

14,145,927 rows (approx. 10%)



full_train_results.csv

13.25 GB

full_test_results.csv

1.66 GB

full_val_results.csv

1.66 GB

Preprocessing Artifact: MinMax Normalizer

- Our feature set includes 7 continuous features with vastly different scales. Training a neural net without first scaling numerical input can lead to erratic or no learning.
- To account for this, we normalize continuous features to a zero to one scale. The scaler must be fit using all 113M rows in the training set, and the same scaler fit on the training set must be saved and used to scale continuous features in the validation and test sets for evaluation and inference respectively.

distance	hour_start	month_start	pickup_latitude	pickup_longitude	dropoff_latitude	dropoff_longitude			
4092.426273	20	10	41.949829	-87.643965	41.919225	-87.671446			
2239.665961	21	1	42.001571	-87.695013	42.009623	-87.670167			
3622.427784	18	11	41.946511	-87.806020	41.954028	-87.763399			
4671.111033	3	1	distance	hour_start	month_start	pickup_latitude	pickup_longitude	dropoff_latitude	dropoff_longitude
2239.665961	1	7	0.7620	0.95	0.9	0.104399	1.000000	0.000000	0.903663
			0.0000	1.00	0.0	0.879398	0.685000	1.000000	0.916232
			0.5687	0.85	1.0	0.054702	0.000000	0.384994	0.000000
			1.0000	0.10	0.0	0.000000	0.546212	0.330046	1.000000
			0.0000	0.00	0.6	1.000000	0.838317	0.910929	0.672062

Preprocessing During Training and Inference

Normalization and one-hot encoding

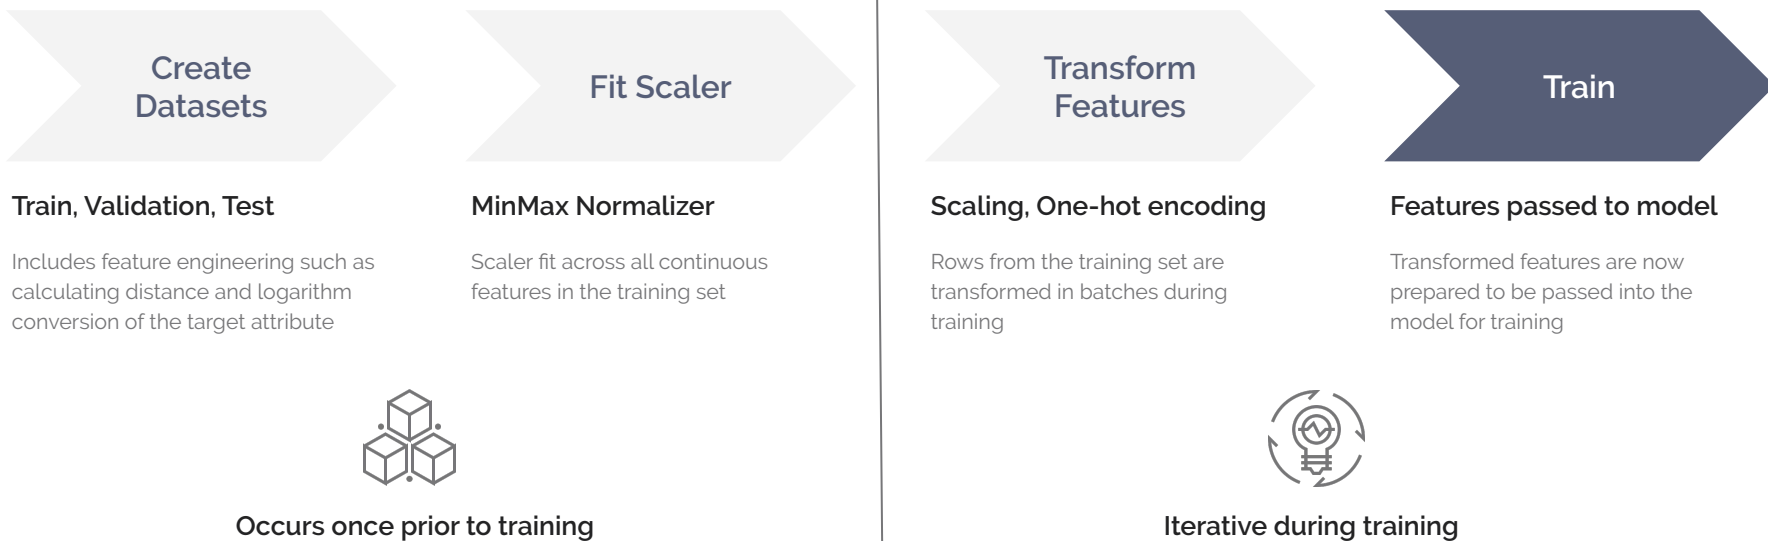
- During training and prior to inference, continuous features are normalized using the MinMax Normalizer artifact
- To be properly fed into the model, categorical features must be transformed using "one-hot" encoding so that each distinct categorical value becomes a column where the rows are either zero or one. It's important that all one-hot encoding reflect the categorical values and order that are available in the *training* set. For example, if a categorical value is found in the validation set that is not found in the training set, it cannot be included in the transformed one-hot encoded matrix.

pickup_census_tract

	9999				
	pickup_census_tract_17031031400	pickup_census_tract_17031062000	pickup_census_tract_17031062800	pickup_census_tract_17031063100	pickup_census_tract_9999
17031031400	0	0	0	0	1
17031063100	1	0	0	0	0
17031062800	0	0	1	0	0
17031062000	0	1	0	0	0

Preprocessing Data Pipeline

One-time and Recurring steps



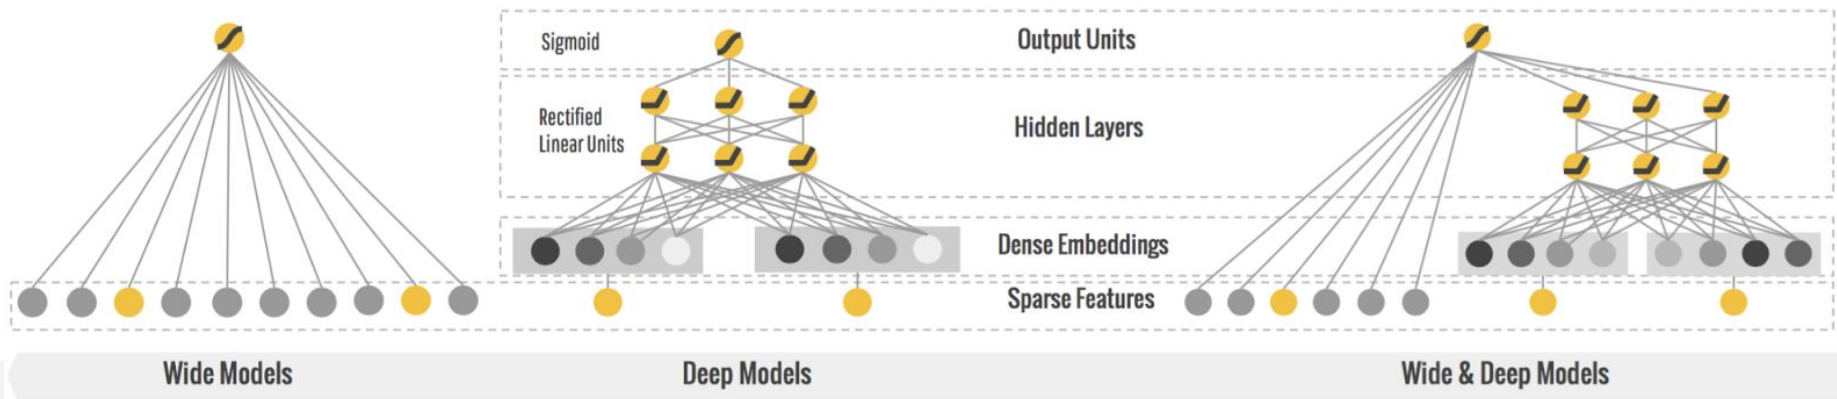
Modeling

Wide and Deep Neural Network

- Wide layers to encode the categorical features
- Deep layers to learn the continuous features
- Architecture optimized with hyperparameter tuning
- Training and Hosting on GCP

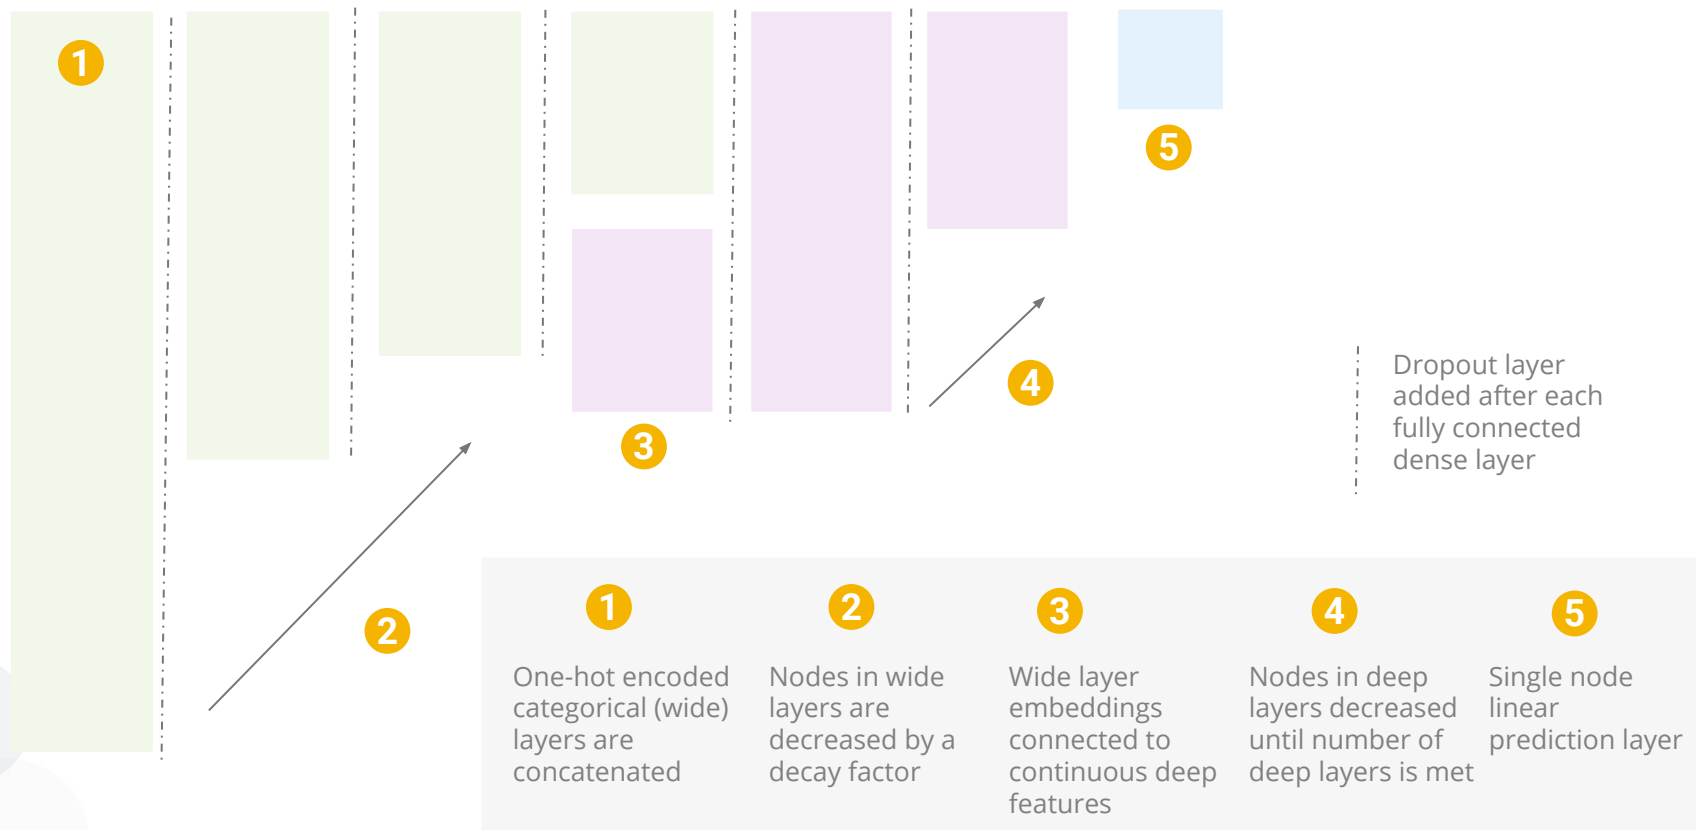
Wide and Deep Neural Networks

- By jointly training a wide linear model (for memorization of categorical features) alongside a deep neural network (for generalization of continuous features), we can combine the strengths of both to make better predictions
- Wide and deep model architectures are useful for generic large-scale regression and classification problems with sparse inputs (categorical features with a large number of possible feature values), such as we have in this dataset with the pickup and dropoff census tract and community area fields.



source: <https://ai.googleblog.com/2016/06/wide-deep-learning-better-together-with.html>

Network Architecture



Hyperparameter Tuning

Optimizing Model Performance

- Our model can be optimized by tuning hyperparameters at the model level such as the learning rate, as well as those that impact the architecture of the model itself such as the number of nodes in the first wide layer and the dropout rate.
- GCP's AI Platform can be used to run trials that iterate through different configurations of the hyperparameters provided with the goal of minimizing or maximizing a target metric - in this case, we minimize the Mean Squared Error on the validation set.

parameter	data type	range	scaling method
train-batch-size	INTEGER	(8, 256)	UNIT_LINEAR_SCALE
learning-rate	DOUBLE	(0.001, 0.05)	UNIT_LOG_SCALE
num-deep-layers	INTEGER	(1, 2, 3, 4)	DISCRETE
first-deep-layer-size	INTEGER	(20, 30, 40, 50)	DISCRETE
first-wide-layer-size	INTEGER	(700, 5000)	UNIT_LINEAR_SCALE
wide-scale-factor	DOUBLE	(0.001, 0.5)	UNIT_LOG_SCALE
dropout-rate	DOUBLE	(0.01, 0.2, 0.3, 0.4, 0.5)	DISCRETE

HyperTune trials

	Trial ID	val_mse ↑	Training step	train-batch-size	learning-rate	num-deep-layers	first-deep-layer-size	first-wide-layer-size
○	31	0.185	40	132	0.003	2	30	1,233
○	39	0.19	40	59	0.024	2	20	2,028
○	13	0.193	40	229	0.023	3	20	3,376
○	43	0.198	49	155	0.007	3	50	2,333
○	5	0.202	37	194	0.014	1	30	3,142
○	9	0.203	49	11	0.005	3	40	4,488
○	28	0.203	49	168	0.004	3	50	1,162

Modeling on Google Cloud Platform

Traceability and Scale with AI Platform

- Training, hyperparameter, and hosting all are completed using AI Platform on Google Cloud Platform
- Access to the models is secured and controlled using IAM roles
- The roles described below define which users have access to which AI Platform resources

Role Title	Role Name	Capabilities
AI Platform Admin	roles/ml.admin	Full control of AI Platform project, and its jobs, operations, models, and versions. Note: The primitive project Editor role is equivalent to roles/ml.admin .
AI Platform Developer	roles/ml.developer	Create training and prediction jobs, models and versions, and send online prediction requests.
AI Platform Viewer	roles/ml.viewer	Read-only access to AI Platform resources.

source: <https://cloud.google.com/ml-engine/docs/access-control>


Evaluation

- Target Variable Transformations
- Evaluation metrics and results
- Suggestions for improvement

Impact of Target Variable Transformation

Training Metrics vs Evaluation Metrics

- As a reminder, we transformed the target attribute (trip duration in seconds) prior to training by taking the natural log of the actual duration to give it a more normal distribution.
- As a result, all metrics generated during training are also associated with the log transform of trip duration
- To assess metrics associated with the actual trip duration in seconds, we'll need to first make predictions for our test set using the trained model



Choosing an Evaluation Metric

Finding the Right Performance Metric for your Use Case

Accuracy

Accuracy is the most commonly considered and well understood evaluation metric - how often are we right? However, accuracy is not a great metric for continuous the target variables like we have in prediction because we won't often be exactly right. Fortunately there are additional metrics can help us understand how close our metrics are to the ground truth.

Root Mean Squared Error (RMSE)

RMSE is a common evaluation metric in prediction problems because it has the benefit of being in the same units as the target variable. For example we might say that, on average, our predictions are within "X" seconds of the actual trip duration. As we've, some trips on our dataset are much longer than others. Being off by 5 seconds on a trip that is 90 minutes is not the same as if the trip were 3 minutes.

Mean Absolute Percentage Error (MAPE)

Fortunately there are metrics that can help us account for these differences in scale. The Mean Absolute Percentage error tells us, on average, by what percentage our predictions were off the mark.

Mean Absolute Percentage Error



Samples

- 5000 trips from the "unseen" test set were used to approximate a MAPE value for our model



Methodology

- Each test sample was passed through preprocessing and then through our trained model so that the prediction could be made and compared against the actual value



Results

- For the subset in question, our predictions were, on average, within **41%** of the actual trip duration in seconds.

Recommendations for Improvement

Optimizing Model Performance is Ongoing

- Additional / third party data sources can be included to add more predictive power to the model. Some examples that may be helpful in this scenario are:
 - Weather data
 - Events data
 - Google Distance Matrix API ;)
- Further exploration with feature selection and feature engineering
 - Does a model work better with only the census tract and not the community area detail?
 - Are there "cross features" - such as a combined field representing the census tract pickup - to - census tract dropoff that get better performance than considering these fields independently?
- Continued hyperparameter and model architecture tuning. Note that the range of hyperparameters included in our hyperparameter search were limited to those we passed in using the configuration. It's likely that there are additional model architectures and / or parameters that would improve model performance.